

1 Logistic Regression: Cost, Gradient and Hessian

Sigmoid Function

$$g(z) = \frac{1}{1 + e^{-z}}.$$

Hypothesis Function $h_{\theta}(x)$

$$h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}.$$

Given a training set $\{(x^{(i)}, y^{(i)})\}_{i=1}^m$ with $y^{(i)} \in \{0, 1\}$ and $x^{(i)} \in \mathbb{R}^n$, the *average empirical loss* (negative log-likelihood) is

Cost Function $J(\theta)$

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m \left[y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \right].$$

1.1 Gradient of $J(\theta)$

First, write down the gradient component-wise. For each training example i , we have

Gradient Derivation (Part 1)

$$h_{\theta}(x^{(i)}) = g(\theta^T x^{(i)}), \quad g'(z) = g(z) (1 - g(z)).$$

Gradient: Chain-Rule Expansion

$$\nabla_{\theta} J(\theta) = -\frac{1}{m} \sum_{i=1}^m \left[y^{(i)} \frac{1}{h_{\theta}(x^{(i)})} \nabla_{\theta} h_{\theta}(x^{(i)}) - (1 - y^{(i)}) \frac{1}{1 - h_{\theta}(x^{(i)})} \nabla_{\theta} h_{\theta}(x^{(i)}) \right].$$

Now, since

$$\nabla_{\theta} h_{\theta}(x^{(i)}) = g'(\theta^T x^{(i)}) x^{(i)} = h_{\theta}(x^{(i)}) (1 - h_{\theta}(x^{(i)})) x^{(i)},$$

substitute into the above:

Gradient: Simplification

$$\begin{aligned}\nabla_{\theta} J(\theta) &= -\frac{1}{m} \sum_{i=1}^m \left[y^{(i)} \frac{1}{h_{\theta}(x^{(i)})} - (1 - y^{(i)}) \frac{1}{1 - h_{\theta}(x^{(i)})} \right] \left[h_{\theta}(x^{(i)}) (1 - h_{\theta}(x^{(i)})) x^{(i)} \right] \\ &= -\frac{1}{m} \sum_{i=1}^m \left[y^{(i)} (1 - h_{\theta}(x^{(i)})) - (1 - y^{(i)}) h_{\theta}(x^{(i)}) \right] x^{(i)} \\ &= -\frac{1}{m} \sum_{i=1}^m \left[y^{(i)} - y^{(i)} h_{\theta}(x^{(i)}) - h_{\theta}(x^{(i)}) + y^{(i)} h_{\theta}(x^{(i)}) \right] x^{(i)} \\ &= -\frac{1}{m} \sum_{i=1}^m \left[y^{(i)} - h_{\theta}(x^{(i)}) \right] x^{(i)}.\end{aligned}$$

Hence, in compact vector form:

Gradient of $J(\theta)$

$$\nabla_{\theta} J(\theta) = -\frac{1}{m} \sum_{i=1}^m (y^{(i)} - h_{\theta}(x^{(i)})) x^{(i)}.$$

1.2 Hessian of $J(\theta)$

Recall the Hessian is $H = \nabla_{\theta}^2 J(\theta)$. We can either differentiate the gradient again, or observe directly that

$$\nabla_{\theta} h_{\theta}(x^{(i)}) = h_{\theta}(x^{(i)}) (1 - h_{\theta}(x^{(i)})) x^{(i)}.$$

When we differentiate $\nabla_{\theta} J(\theta)$ another time, each term brings down another copy of $x^{(i)}$. Concretely:

Hessian of $J(\theta)$

$$H = \nabla_{\theta}^2 J(\theta) = \frac{1}{m} \sum_{i=1}^m h_{\theta}(x^{(i)}) (1 - h_{\theta}(x^{(i)})) x^{(i)} (x^{(i)})^T.$$

—

1.3 Proof that H is Positive Semi-Definite

For any $z \in \mathbb{R}^n$, consider

Positive Semi-Definiteness

$$\begin{aligned} z^T H z &= z^T \left[\frac{1}{m} \sum_{i=1}^m h_{\theta}(x^{(i)}) (1 - h_{\theta}(x^{(i)})) x^{(i)} (x^{(i)})^T \right] z \\ &= \frac{1}{m} \sum_{i=1}^m h_{\theta}(x^{(i)}) (1 - h_{\theta}(x^{(i)})) (z^T x^{(i)}) ((x^{(i)})^T z) \\ &= \frac{1}{m} \sum_{i=1}^m h_{\theta}(x^{(i)}) (1 - h_{\theta}(x^{(i)})) (z^T x^{(i)})^2 \geq 0, \end{aligned}$$

showing H is positive semi-definite.

ps1__b

June 6, 2025

```
[3]: import numpy as np
import util

from linear_model import LinearModel
import matplotlib.pyplot as plt
from sklearn.metrics import classification_report, confusion_matrix

class LogisticRegression(LinearModel):
    @staticmethod
    def sigmoid(z):
        return np.where(z >= 0, 1 / (1 + np.exp(-z)), np.exp(z) / (1 + np.
↪exp(z)))

    @staticmethod
    def compute_htheta(X, theta):
        return LogisticRegression.sigmoid(np.dot(X, theta))

    @staticmethod
    def compute_grad(X,y,theta):
        m = X.shape[0]
        h = LogisticRegression.compute_htheta(X, theta)           # shape (m,)
        error = y - h                                             # shape (m,)
        grad = -(X.T @ error) / m
        return grad

    @staticmethod
    def compute_hessian(X,theta):
        hessian = np.zeros((X.shape[1], X.shape[1]))
        m = X.shape[0]
        h_theta = LogisticRegression.compute_htheta(X, theta)
        for i in range(m):
            hessian += h_theta[i] * (1 - h_theta[i]) * np.outer(X[i], X[i])
        hessian /= m
        return hessian

    def fit(self, x, y):
        # *** START CODE HERE ***
```

```

theta = np.zeros(x.shape[1])

max_iter = 1000
for _ in range(max_iter):
    gradient_matrix = LogisticRegression.compute_grad(x, y, theta)
    hessian_matrix = LogisticRegression.compute_hessian(x, theta)

    delta = np.linalg.solve(hessian_matrix, gradient_matrix)
    if np.linalg.norm(delta) < 1e-6:
        break
    theta -= delta

self.theta = theta.copy()

# *** END CODE HERE ***

def predict(self, x):
    """Make a prediction given new inputs x.

    Args:
        x: Inputs of shape (m, n).

    Returns:
        Outputs of shape (m,).
    """
    # *** START CODE HERE ***
    probs = LogisticRegression.compute_htheta(x, self.theta)
    return (probs >= 0.5).astype(int)
    # *** END CODE HERE ***

def main(train_path, eval_path, pred_path):
    """Problem 1(b): Logistic regression with Newton's Method.

    Args:
        train_path: Path to CSV file containing dataset for training.
        eval_path: Path to CSV file containing dataset for evaluation.
        pred_path: Path to save predictions.
    """
    x_train, y_train = util.load_dataset(train_path, add_intercept=True)
    x_eval, y_eval = util.load_dataset(eval_path, add_intercept=True)

    # *** START CODE HERE ***
    clf = LogisticRegression()
    clf.fit(x_train, y_train)

    y_pred = clf.predict(x_eval)

```

```

np.savetxt(pred_path, y_pred, fmt="%d")

acc = np.mean(y_pred == y_eval)
print(f"Eval accuracy = {acc:.4f}")

print("Confusion Matrix:")
print(confusion_matrix(y_eval, y_pred))

print("\nClassification Report:")
print(classification_report(y_eval, y_pred))

w0, w1, w2 = clf.theta

x1_vals = np.linspace(min(x_eval[:, 1]), max(x_eval[:, 1]), 100)
x2_vals = -(w0 + w1 * x1_vals) / w2

plt.scatter(x_eval[:, 1], x_eval[:, 2], c=y_eval, cmap="bwr", alpha=0.7)
plt.plot(x1_vals, x2_vals, color="k", linewidth=2)
plt.xlabel("Feature 1")
plt.ylabel("Feature 2")
plt.title("Decision Boundary on Validation Set")
plt.show()

# *** END CODE HERE ***
if __name__ == "__main__":
    import sys
    main("./data/ds1_train.csv", "./data/ds1_valid.csv", "./data/ds1_pred.csv")

```


Eval accuracy = 0.9000

Confusion Matrix:

```
[[42  8]
 [ 2 48]]
```

Classification Report:

	precision	recall	f1-score	support
0.0	0.95	0.84	0.89	50
1.0	0.86	0.96	0.91	50
accuracy			0.90	100
macro avg	0.91	0.90	0.90	100
weighted avg	0.91	0.90	0.90	100



2 Gaussian Discriminant Analysis: Closed-Form for $p(y = 1 | x)$

In GDA, we assume

$$y \sim \text{Bernoulli}(\phi), \quad p(x | y = 0) = \mathcal{N}(x; \mu_0, \Sigma), \quad p(x | y = 1) = \mathcal{N}(x; \mu_1, \Sigma),$$

where $\phi \in (0, 1)$, $\mu_0, \mu_1 \in \mathbb{R}^n$, and $\Sigma \in \mathbb{R}^{n \times n}$ is positive-definite and common to both classes. Then by Bayes' rule:

Posterior Probability $p(y = 1 | x)$

$$\begin{aligned} p(y = 1 | x; \phi, \mu_0, \mu_1, \Sigma) &= \frac{p(x | y = 1; \mu_1, \Sigma) p(y = 1; \phi)}{p(x | y = 1; \mu_1, \Sigma) p(y = 1; \phi) + p(x | y = 0; \mu_0, \Sigma) p(y = 0; \phi)} \\ &= \frac{1}{1 + \exp\left(\log\left[\frac{p(x|y=0) p(y=0)}{p(x|y=1) p(y=1)}\right]\right)}. \end{aligned}$$

Because $p(x | y = 1)$ and $p(x | y = 0)$ share the same covariance Σ , we can rewrite the denominator and get a sigmoid form. In particular:

Sigmoid Form of Posterior

$$p(y = 1 | x) = \frac{1}{1 + \frac{p(x|y=0) p(y=0)}{p(x|y=1) p(y=1)}} = \frac{1}{1 + \exp\left(\log\left[\frac{p(x|y=0) p(y=0)}{p(x|y=1) p(y=1)}\right]\right)}.$$

For a multivariate normal,

Gaussian Density $p(x | y = j)$

$$p(x | y = j) = \frac{1}{(2\pi)^{n/2} |\Sigma|^{1/2}} \exp\left\{-\frac{1}{2} (x - \mu_j)^T \Sigma^{-1} (x - \mu_j)\right\}, \quad j = 0, 1.$$

Therefore, the log-ratio becomes

Log-Ratio of Class-Conditional Terms

$$\log \frac{p(x | y = 0) p(y = 0)}{p(x | y = 1) p(y = 1)} = -\frac{1}{2} (x - \mu_0)^T \Sigma^{-1} (x - \mu_0) + \frac{1}{2} (x - \mu_1)^T \Sigma^{-1} (x - \mu_1) + \log \frac{1 - \phi}{\phi}.$$

Plugging back into the sigmoid form:

Posterior in Terms of Exponential

$$p(y = 1 | x) = \frac{1}{1 + \exp \left(\underbrace{-\frac{1}{2} (x - \mu_1)^T \Sigma^{-1} (x - \mu_1) + \frac{1}{2} (x - \mu_0)^T \Sigma^{-1} (x - \mu_0) + \log \frac{1 - \phi}{\phi}}_{= -(\theta^T x + \theta_0)} \right)}.$$

Re-arrange the quadratic forms:

Quadratic Expansion

$$\begin{aligned} & -\frac{1}{2} (x - \mu_1)^T \Sigma^{-1} (x - \mu_1) + \frac{1}{2} (x - \mu_0)^T \Sigma^{-1} (x - \mu_0) + \log \frac{1 - \phi}{\phi} \\ &= -\frac{1}{2} \left[x^T \Sigma^{-1} x - 2 \mu_1^T \Sigma^{-1} x + \mu_1^T \Sigma^{-1} \mu_1 \right] + \frac{1}{2} \left[x^T \Sigma^{-1} x - 2 \mu_0^T \Sigma^{-1} x + \mu_0^T \Sigma^{-1} \mu_0 \right] + \log \frac{1 - \phi}{\phi} \\ &= -x^T \Sigma^{-1} x + \mu_1^T \Sigma^{-1} x - \frac{1}{2} \mu_1^T \Sigma^{-1} \mu_1 + \frac{1}{2} x^T \Sigma^{-1} x - \mu_0^T \Sigma^{-1} x + \frac{1}{2} \mu_0^T \Sigma^{-1} \mu_0 + \log \frac{1 - \phi}{\phi} \\ &= -(\mu_1 - \mu_0)^T \Sigma^{-1} x - \left[\frac{1}{2} \mu_1^T \Sigma^{-1} \mu_1 - \frac{1}{2} \mu_0^T \Sigma^{-1} \mu_0 \right] + \log \frac{1 - \phi}{\phi}. \end{aligned}$$

Define

GDA Parameter Definitions

$$\theta = \Sigma^{-1} (\mu_1 - \mu_0), \quad \theta_0 = -\frac{1}{2} \mu_1^T \Sigma^{-1} \mu_1 + \frac{1}{2} \mu_0^T \Sigma^{-1} \mu_0 - \log \left(\frac{1 - \phi}{\phi} \right).$$

Then

Final Posterior Form

$$p(y = 1 | x; \phi, \mu_0, \mu_1, \Sigma) = \frac{1}{1 + \exp(-(\theta^T x + \theta_0))},$$
$$\theta = \Sigma^{-1} (\mu_1 - \mu_0), \quad \theta_0 = -\frac{1}{2} \mu_1^T \Sigma^{-1} \mu_1 + \frac{1}{2} \mu_0^T \Sigma^{-1} \mu_0 - \log \left(\frac{1 - \phi}{\phi} \right).$$

gda

June 6, 2025

```
[5]: import numpy as np
import util

from linear_model import LinearModel
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt

def plot_decision_boundary(x, y, theta):
    plt.figure()

    # Plot data points
    plt.scatter(x[y==0][:, 0], x[y==0][:, 1], label="Class 0", marker='o')
    plt.scatter(x[y==1][:, 0], x[y==1][:, 1], label="Class 1", marker='x')

    # Plot decision boundary:  $w_0 + x_1 + x_2 = 0 \implies x_2 = -(w_0 + x_1) / w_2$ 
    x1_vals = np.linspace(np.min(x[:,0]), np.max(x[:,0]), 100)
    theta_0, theta_1, theta_2 = theta
    x2_vals = -(theta_0 + theta_1 * x1_vals) / theta_2
    plt.plot(x1_vals, x2_vals, label="Decision Boundary", color='green')

    plt.xlabel("x1")
    plt.ylabel("x2")
    plt.legend()
    plt.title("GDA Decision Boundary")
    plt.grid(True)
    plt.show()

class GDA(LinearModel):
    """Gaussian Discriminant Analysis.

    Example usage:
    > clf = GDA()
    > clf.fit(x_train, y_train)
    > clf.predict(x_eval)
    """
```

```

def fit(self, x, y):
    """Fit a GDA model to training set given by x and y.

    Args:
        x: Training example inputs. Shape (m, n).
        y: Training example labels. Shape (m,).

    Returns:
        theta: GDA model parameters.
    """
    # *** START CODE HERE ***
    phi = np.mean(y)
    mu_0 = np.mean(x[y == 0], axis=0)
    mu_1 = np.mean(x[y == 1], axis=0)
    cov_matrix = np.zeros((x.shape[1], x.shape[1]))
    for i in range(x.shape[0]):
        if y[i] == 0:
            cov_matrix += np.outer(x[i] - mu_0, x[i] - mu_0)
        else:
            cov_matrix += np.outer(x[i] - mu_1, x[i] - mu_1)
    cov_matrix /= x.shape[0]
    inverse_cov = np.linalg.inv(cov_matrix)

    theta_0 = -((mu_1.T @ inverse_cov @ mu_1)/2 + (mu_0.T @ inverse_cov @
↪mu_0)/2 + np.log(phi / (1 - phi)))

    theta = inverse_cov @ (mu_1 - mu_0)

    theta = np.concatenate([theta_0, theta])
    self.theta = theta

    # *** END CODE HERE ***

def predict(self, x):
    """Make a prediction given new inputs x.

    Args:
        x: Inputs of shape (m, n).

    Returns:
        Outputs of shape (m,).
    """
    # *** START CODE HERE ***
    if self.theta is None:

```

```

        raise ValueError("Model has not been fitted yet. Call fit() before_
predict().")
    x_new = np.concatenate([np.ones((x.shape[0], 1)), x], axis=1)
    return x_new @ self.theta >= 0

    # *** END CODE HERE
def main(train_path, eval_path, pred_path):
    # Load dataset
    x_train, y_train = util.load_dataset(train_path, add_intercept=False)

    # *** START CODE HERE ***
    x_eval, y_eval = util.load_dataset(eval_path, add_intercept=False)

    clf = GDA()
    clf.fit(x_train, y_train)
    y_pred = clf.predict(x_eval)
    np.savetxt(pred_path, y_pred)

    acc = accuracy_score(y_eval, y_pred)
    print(f"Accuracy on evaluation set: {acc:.4f}")

    if x_eval.shape[1] == 2:
        plot_decision_boundary(x_eval, y_eval, clf.theta)

    # *** END CODE HERE ***

if __name__ == "__main__":
    main("data/ds1_train.csv", "data/ds1_valid.csv", "output/predictions.txt")

```

Accuracy on evaluation set: 0.8300

